Judit — Operator workflows

Run, inspect, repair, review.

Source-of-truth: `docs/dev/`, `docs/reference/`, `README.md`.

What operators actually do

1. Curate the **source registry** — register, refresh, and select legal sources for runs.
2. Run **proposition analysis** — single-jurisdiction inventory or comparison/divergence mode.
3. Inspect **runs and artifacts** — quality summary, extraction traces, repair signals.
4. Triage and **review** — record decisions across review layers.
5. Repair or rerun where needed; promote to **guidance-ready** under governance policy.

Source registry and discovery

- CLI: `source-registry-list`, `source-registry-inspect`, `source-registry-register`, `source-registry-refresh`, `run-registry-sources`.
- API: `GET /ops/source-registry`, `POST /ops/source-registry/register`, `POST /ops/source-registry/search`, `POST /ops/source-registry/{registry_id}/refresh`.
- Run from registry: `POST /ops/runs/from-registry` (sync) · `POST /ops/run-jobs/from-registry` (async, preferred for the operator UI).

Source intake layers

- **Registered sources** — persistent registry entries, selectable for runs.
- **Source records in a run** — written to `sources.json` after ingestion.
- **Source-family candidates** — `source_family_candidates.json` — not analysed until registered.
- **Attached context** — audit/review metadata; not legal source input unless promoted.

See `docs/dev/source-registry.md`.

Fragment selection modes

Fragment selection modes

Three policies trade reviewability, exploration, and cost. Every selection is auditable.

`required_only`

Reviewable / staged profiles

- Only fragments matching `required_fragment_locators` run.
- Explicit annex inclusion only; no focus-term auto-selection.
- Best for staged reviewable profiles where each fragment must be triable by a reviewer.
- Lowest cost / noise. Tight scope - confirms what was deliberately registered.

`required_plus_focus`

Exploratory profile

- Required locators always run.
- Adds fragments whose text matches configured focus terms (e.g., scope keywords).
- Broader exploratory coverage without going wide on every match in the corpus.
- Moderate cost / noise; useful when extending a staged profile.

`all_matching`

Broad / exploratory

- Broad matching - close to the historical default behaviour.
- Useful for early exploration of an unknown corpus or sweep runs.
- Highest cost / noise risk. Pair with strong quality gates and triage budget.
- Not recommended as the default for staged reviewable profiles.

Selection audit (always written to proposition_extraction_jobs.json)

`selected_for_extraction`
Whether the fragment was actually selected to run extraction.

`skip_reason`
Why a candidate fragment was skipped under the active policy.

`estimated_input_tokens`
Conservative chunk size estimate used for context-window safety.

`selection_reason`
Why the fragment was chosen (required, focus match, broad match, etc.).

`extraction_mode & model_alias`
heuristic / local / frontier mode and the model alias used when applicable.

`repairable, repair_reason`
Marks rows eligible for the repair workflow (not legal acceptance).

Selection controls cost, noise, and reviewability. Audit fields make every run inspectable, regardless of mode.

Extraction outcomes & repair

Extraction outcomes & repair

Fragment-to-validation flow: clean, fallback, repairable, and fail-closed outcomes - all reviewable.

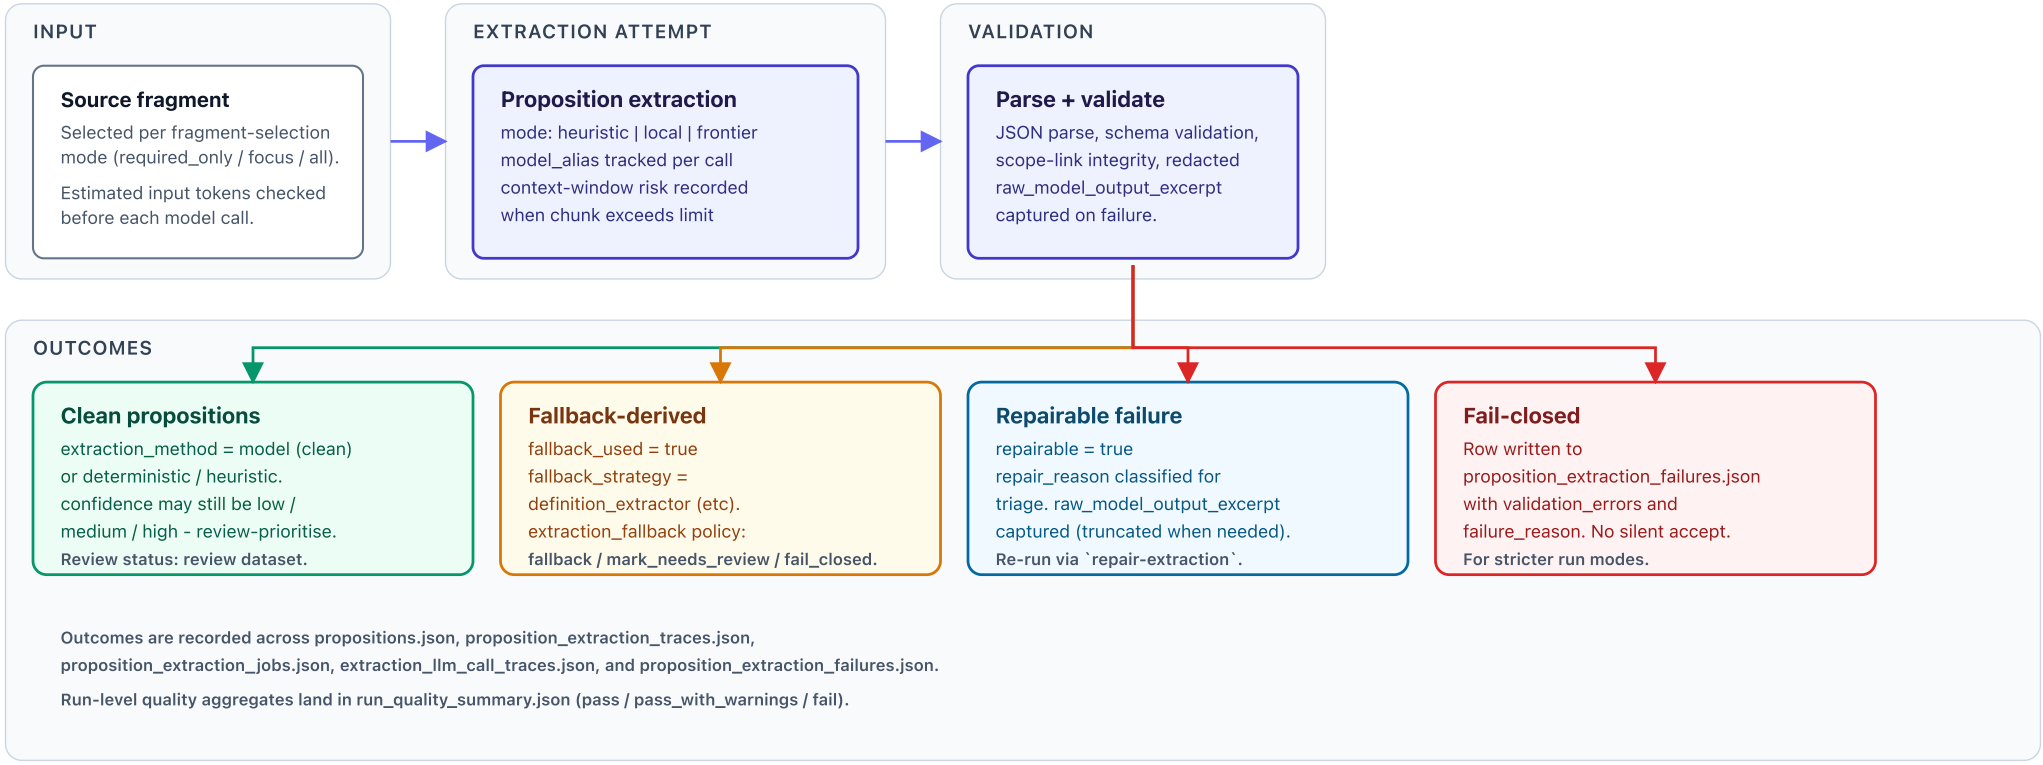

Outcomes are recorded across propositions.json, proposition_extraction_traces.json, proposition_extraction_jobs.json, extraction_llm_call_traces.json, and proposition_extraction_failures.json.

Run-level quality aggregates land in run_quality_summary.json (pass / pass_with_warnings / fail).

REVIEWABLE

Operator inspects repairable rows, fallback traces, low-confidence traces, and parse-error excerpts via the operations UI / `/ops/*` endpoints. Reviewers record `pipeline_review_decisions` (approved / rejected /

Run quality

What "run quality" means (and doesn't)

run_quality_summary.json is an artifact-integrity gate. It is not legal approval.

`pass`

All quality gates satisfied with no warnings flagged.

Indicates artifact integrity and gate compliance, not that propositions are legally correct.

Review may still be required by governance policy before promotion to guidance-ready.

`pass_with_warnings`

Gates satisfied, but actionable warnings are present.

Expected during extraction/review work - should be triaged, not ignored.

Common drivers: fallback usage, low-confidence traces, context-window risks, scope-link noise.

`fail`

One or more gates failed.

Treat as a halt for downstream consumption.

Drivers include validation_errors, fail-closed extraction failures, or critical lint breakages. Inspect proposition_extraction_failures.json.

What a lint pass means

- ✓ All expected artifacts are present and well-formed.
- ✓ Cross-references between artifacts are consistent.
- ✓ Schemas validate; required fields are populated.
- ✓ Operational gates (counts, ratios, repairable share) within thresholds.
- ✓ Run is safe to inspect, hand off, and review.

Inspect via:

```
uv run --package judit-pipeline python -m judit_pipeline \
lint-export --export-dir dist/static-report
```

What a lint pass does not mean

- ✕ Propositions are legally correct or court-ready.
- ✕ Coverage of the legal area is complete.
- ✕ Promotion to guidance-ready is automatic.
- ✕ Fallback / repairable / low-confidence rows are accepted.
- ✕ Human legal review can be skipped.
- ✕ A model-assisted stage produced a deterministic answer.

Decisions, overrides, and supersession are recorded in pipeline_review_decisions.json - that is the governance trail.

PROMOTION PATH

run → quality pass / pass_with_warnings → reviewer decisions across layers (proposition, structured display, scope links, completeness) → extraction confidence is high · completeness not fragmentary · no unresolved repairable failures → guidance-ready candidacy →

Review governance

Review governance workflow

Append-only decisions across review layers - promotion to guidance-ready is earned, not default.

REVIEW LAYERS (PER PROPOSITION ROW)

proposition
The atomic legal statement extracted from a fragment.

structured_proposition_display
Subject / rule / object / condition view shown in the UI.

proposition_scope_links
Links to scope taxonomy (equine, bovine, etc.) - quieter by default.

proposition_completeness
Self-contained vs fragmentary; context-aware display.

proposition_extraction_trace · source_target_link · source_categorisation_rationale · source_inver
Lineage and structure rows are reviewable too: confidence, role assignments, target links, and source-inventory entries can each carry a decision.

DECISION VALUES (PIPELINE_REVIEW_DECISIONS.JSON)

approved	Accepted under governance policy.
rejected	Not acceptable as-is; remediation or drop required.
needs_review	Held for further reviewer attention.
overridden	Field-scoped replacement_value applied with rationale.
deferred	Acknowledged, decision postponed with explicit reason.

Decisions are run-scoped, append-only, and may carry `supersedes_decision_id` to retire prior rows. Effective decision is the latest non-superseded row for the artifact (+ optional field).

Mutation: POST /ops/runs/{run_id}/pipeline-review-decisions · CLI: add-review-decision.

DECISION PIPELINE

Review decision
Reviewer chooses a value for the artifact + optional applies_to_field. Reason and evidence may be recorded.

→

Append-only record
New row written with unique id, reviewer, reviewed_at, run_id, artifact_type/id linkage.

→

Supersession (opt.)
If a prior decision was superseded, the new row sets supersedes_decision_id - no edits, only replacements.

→

Run/export reload
Decisions are persisted in export artifacts and surface on next read via /ops/pipeline-review-decisions.

→

guidance-ready candidacy
Currently most explicit in equine coverage: accepted (or accepted_with_edits) · extraction confidence high · completeness not fragmentary · scope-links accepted or deferred · no unresolved repairable failure affecting the proposition.

Equine staged corpus

Equine staged corpus roadmap

Three staged profiles for interactive build and review. Coverage is staged - not a complete-corpus claim.

`equine_passport_eu_2015_262_v0_1`

Baseline · EU equine passport

Purpose: EU 2015/262 equine passport baseline.

Corpus config
`examples/corpus_equine_passport_eu_2015_262_v0_1.json`

Case
`examples/equine_passport_eu_2015_262_case.json`

Estimate
```
estimate-corpus-run --profile equine_passport_eu_2015_262_v0_1 --fe
```

Build (requires LiteLLM + provider credit)
```
build-equine-corpus \
--corpus-config examples/corpus_equine_passport_eu_2015_262_v0_1.js
--output-dir dist/static-report --use-llm
```

`equine_passport_england_ukwide_v0_1`

England + UK-wide retained / EU Exit

Purpose: England + UK-wide retained / EU Exit equine identification layer.

Corpus config
`examples/corpus_equine_passport_england_ukwide_v0_1.json`

Case
`examples/equine_passport_england_ukwide_case.json`

Estimate
```
estimate-corpus-run --profile equine_passport_england_ukwide_v0_1 -
```

Build
```
build-equine-corpus \
--corpus-config examples/corpus_equine_passport_england_ukwide_v0_1
--output-dir dist/static-report --use-llm
```

`equine_passport_devolved_v0_1`

Devolved nations

Purpose: Wales / Scotland / Northern Ireland equine identification layer.

Corpus config
`examples/corpus_equine_passport_devolved_v0_1.json`

Case
`examples/equine_passport_devolved_case.json`

Estimate
```
estimate-corpus-run --profile equine_passport_devolved_v0_1 --fetch
```

Build
```
build-equine-corpus \
--corpus-config examples/corpus_equine_passport_devolved_v0_1.json \
--output-dir dist/static-report --use-llm
```

Umbrella profile: equine_passport_identification_v0_1

Larger, batch-oriented profile. Use staged profiles above for interactive frontier extraction and reviewable runs. Do not run the full 21-source universe interactively unless intentionally testing scale.

After every build run `lint-export` and inspect the run-quality summary, repair banner, fallback counts, low-confidence traces, and context-window risks.

Review & promotion

guidance_ready in equine proposition coverage remains 0 until reviewed. Promotion requires accepted review decisions across proposition, structured display, scope links, and completeness, plus high extraction confidence and resolved repairable failures.

Coverage artifacts: equine_source_coverage.json, equine_proposition_coverage.json, equine_corpus_readiness.json - review datasets, not legal sign-off.

Common workflows (CLI)

```
# Run from a case and export (heuristic, provider-free)
uv run --package judit-pipeline python -m judit_pipeline export-case \
  data/demo/example_case.json \
  --output-dir dist/static-report-tmp \
  --extraction-mode heuristic \
  --extraction-fallback mark_needs_review

# Run with model-backed extraction (requires LiteLLM and credit)
uv run --package judit-pipeline python -m judit_pipeline export-case \
  data/demo/example_case.json \
  --output-dir dist/static-report \
  --extraction-mode local
```

Inspecting a run

```
# 1) extraction job counts
jq '{
  total: length,
  selected: ([.[] | select(.selected_for_extraction == true)] | length),
  repairable: ([.[] | select(.repairable == true)] | length)
}' dist/static-report/proposition_extraction_jobs.json

# 2) quality/lint summary
uv run --package judit-pipeline python -m judit_pipeline lint-export \
  --export-dir dist/static-report

# 3) low-confidence extraction traces
jq '[.[] | select(.confidence == "low")] | length' \
  dist/static-report/proposition_extraction_traces.json

# 4) repairable extraction job rows
jq '[.[] | select(.repairable == true)]' \
  dist/static-report/proposition_extraction_jobs.json
```

Dev reset (read the docs first)

- `clear-operations-runs` removes exported runs and keeps registry.
- `clear-operations-all` removes exported runs and clears registry + caches.

Always run with `--dry-run` first.

See `docs/dev/operations-state-reset.md`.

Operational guardrails

- Lint pass is artifact integrity, not legal approval.
- `pass_with_warnings` is expected and should be triaged.
- Fallback / repairable / low-confidence rows require human review.
- `guidance-ready` is a promotion state after review/governance, not default output.
- Source family discovery is evolving and not a complete legal universe by default.

References

- Operator reference: `docs/reference/cli.md`, `docs/reference/api-ops.md`, `docs/reference/artifacts.md`
- Dev guides: `docs/dev/source-registry.md`, `docs/dev/extraction-repair.md`, `docs/dev/review-governance.md`, `docs/dev/quality-runs.md`, `docs/dev/equine-passport-staged-runs.md`, `docs/dev/operations-state-reset.md`
- Roadmap: `docs/canonical/roadmap.md`