Judit

Source-first, proposition-first legal analysis workbench

Stakeholder overview deck. Source-of-truth: `docs/canonical/`, `docs/architecture/`, `docs/dev/`.

What Judit is

Judit turns legal text into reviewable, traceable analysis.

- **Source-first** — every analysis is grounded in registered legal sources with versioned snapshots and fragment-level lineage.
- **Proposition-first** — atomic legal statements are the unit of analysis, review, and (optionally) comparison.
- **Reviewable** — outputs are review datasets until humans accept them under governance policy.
- **Auditable** — runs persist source records, snapshots, fragments, propositions, traces, jobs, and quality summaries.

Why this shape works

- A clear evidence chain from authoritative source to proposition, with provenance retained at every step.
- Useful outputs without forcing a comparison/divergence workflow.
- Structured comparison units (propositions) when divergence is needed.
- Local/cloud model routing via LiteLLM, with traceable model usage only in explicitly defined stages.
- Future-ready static rendering boundary via the export bundle.

Core workflow

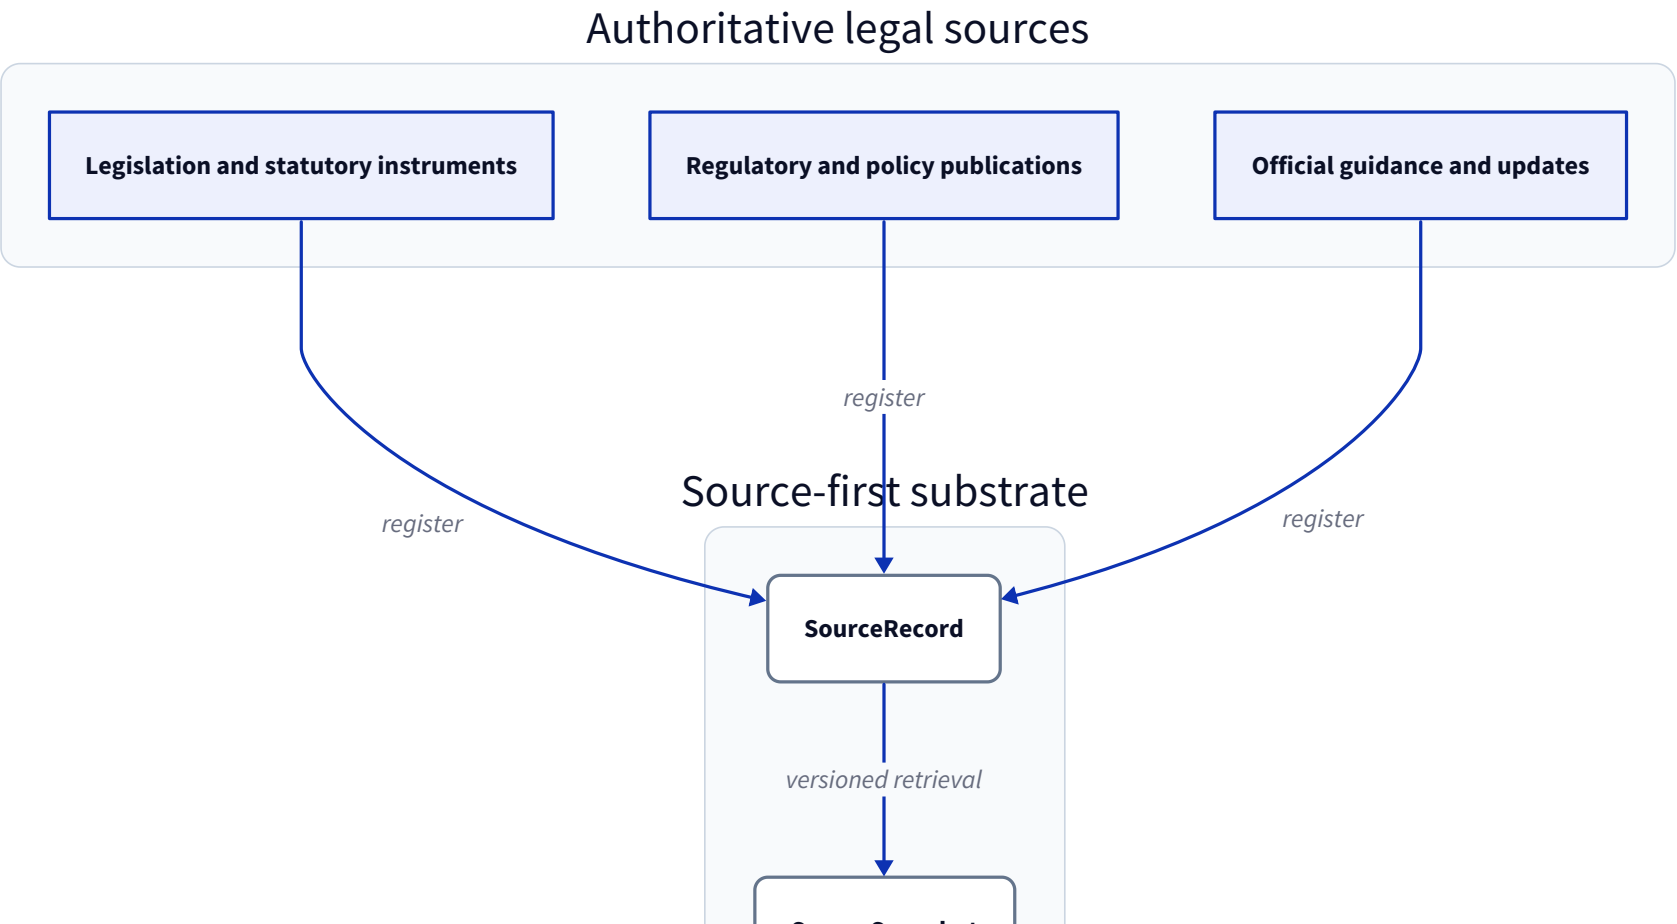

Pipeline at a glance

From legal sources to reviewable propositions

Judit: source-first, proposition-first, deterministic-first pipeline with controlled model usage.

SOURCE INTAKE (NO MODEL)

Source registry
Persisted, selectable references for runs

→ **Source record**
Authority, citation, CELEX/ELI, locator

→ **Source snapshot**
Versioned text instance, content hash, provenance

→ **Source fragment**
Article / annex / section / chunk for extraction targeting

PROPOSITION WORKFLOW

Proposition extraction
Local / frontier model (controlled)
heuristic / local / frontier modes
context-window budgeting
fallback / mark_needs_review
definition fallback path

→ **Inventory & pairing**
Deterministic, no model
group propositions by article / provision
align across sources when divergence is active

→ **Completeness & categorisation**
Deterministic rules
flag fragmentary or context-dependent propositions; assign source roles and relationships

→ **Quality & lint gates**
Deterministic
pass / pass_with_warnings / fail
aggregated metrics + recommendations into run_quality_summary.json

REVIEW & DOWNSTREAM (HUMAN-IN-THE-LOOP)

Operations & analysis UI
Propositions screen, ops inspector, run quality, repair signals, traces. Read-mostly with controlled write paths (registry, decisions, repair).

→ **Pipeline review decisions**
Append-only governance rows per artifact / field: approved, rejected, needs_review, overridden, deferred. Supersession.

→ **Guidance-ready promotion**
Promoted state after review. Not a default output. Not legal approval. Recorded against explicit governance policy.

→ **Divergence (optional)**
Frontier model, traceable
Pair and classify when divergence mode is enabled or multi-jurisdiction input is present.

Lint pass means artifact integrity / quality-gate pass - not legal approval.
`pass_with_warnings` is expected during extraction/review and should be triaged.

`guidance-ready` is a promotion state after review/governance, not default output.
Extracted propositions are review datasets, not legally approved conclusions by default.

Pipeline stages and model boundaries

Stage	What it does	Model usage
Source intake	Fetch, parse, and register sources	No model
Proposition extraction	Extract candidate legal propositions	Local model (controlled, cached)
Inventory & pairing	Group and align propositions	No model
Completeness	Assess whether propositions are self-contained	Deterministic rules

What a proposition looks like

Example

Article 109 §1(d)(i)

Structured proposition

- Subject: Member States
- Rule: must record
- Object: equine identification data
- Condition: within Article 109 system
- Scope: equine

What v1 delivers

- Source registry as the explicit control plane for runs.
- Single-jurisdiction proposition inventory as a first-class workflow.
- Divergence as a downstream mode (enabled explicitly or when multi-jurisdiction input is present).
- Operations and history APIs for runs, sources, traces, and decisions.
- Read-mostly analysis workbench (`/`) and operations/registry inspector (`/ops`).
- LiteLLM gateway and export bundle boundary.

What v1.5 consolidates

- Improved structured proposition reliability and display consistency.
- Provision type classification: core rule, definition, exception, transitional, cross-reference.
- Quieter scope links by default (reviewer can expand).
- Review decisions across proposition, structured display, scope links, and completeness.
- Article/provision grouping as the primary navigation pattern.

Status snapshot (2026-05): structured proposition reliability, provision type surfacing, broader review decisions, and

What v2 and beyond add

- Article-level proposition clusters, related provisions, related sources.
- Cross-reference navigation; comparison across jurisdictions and time.
- Timeline of legal change, "what applies now / later", temporal reasoning across versions.

Explicitly deferred: timeline UI until structured propositions, provision types, and article-level grouping are reliable; LangSmith evaluation; Astro renderer; Slidev; advanced orchestration.

Not a black box

- Deterministic-first pipeline with controlled model usage.
- Models only in explicit stages; every assisted step is traceable.
- Full traceability to sources, snapshots, fragments, and pipeline stages.
- Human review is first-class governance.
- `guidance-ready` is a promotion state earned through review — not default output.
- Lint pass means artifact integrity / quality-gate pass — not legal approval.

Where to go next

- Architecture: `docs/architecture/pipeline-overview.md`, `docs/architecture/system-overview.md`
- Reference: `docs/reference/artifacts.md`, `docs/reference/api-ops.md`, `docs/reference/cli.md`
- Roadmap: `docs/canonical/roadmap.md`, `docs/roadmap/v1.md`
- Decisions: `docs/decisions/adr-0019-model-usage-strategy.md` (model usage), `adr-0018-proposition-identity-and-naming.md` (identity)